DEEP LEARNING · COMPUTER VISION · FINANCIAL TIME SERIES

Market Regime Detection via CNN on Rolling Cross-Asset Correlation Heatmaps

9-asset Pearson correlation matrix → 64 × 64 PNG image → Custom CNN → Bull / Bear / Neutral

0.46	61%	102k	4,878	5.9×
MACRO-F1	BEAR RECALL	PARAMS	HEATMAPS	BEAR LIFT

TensorFlow / Keras · yfinance · seaborn · Weights & Biases · Kaggle T4 GPU · Grad-CAM

MOTIVATION

During market stress, cross-asset correlations **collapse toward +1** — equities co-move while bonds and gold decouple. This structural shift is **invisible to price-level models** yet visible as a distinct spatial pattern in the 30-day Pearson correlation matrix.

We render each matrix as a **64 × 64 heatmap image** and train a CNN to classify regime from pixel structure alone — a pure vision approach to a structural finance problem.

ASSETS & LABELS

EQUITY	MACRO
XLF · XLK · XLE	GLD · TLT
XLV · XLI · VNQ	^VIX (level only)

Regime	Rule
Bear	`VIX>30` OR `drawdown>15%`
Bull	`VIX<20` AND `drawdown<5%`
Neutral	Residual — everything else

9.8% · 61.7% · 28.4%

Bear 483 · Bull 3,028 · Neutral 1,396

CNN ARCHITECTURE

- `Input 64×64×3` — RGB heatmap PNG
- ↓
- `Conv2D(32) + BN` · `MaxPool → 32²`
- `Conv2D(64) + BN` · `MaxPool → 16²`
- `Conv2D(128) + BN` · `MaxPool → 8²`
- `GlobalAvgPool → 128d` — replaces Flatten
- `Dense(64) + Dropout(0.4)`
- `Softmax(3)` — bear / bull / neutral

ARCHITECTURE DESIGN DECISIONS

WHY GAP?
GlobalAvgPool prevents a 524k-param head on only ~3,400 training samples. Spatially averages feature maps instead of flattening.

WHY BATCH NORM?
Normalises activations per-batch, stabilises gradient flow through 3 conv layers, and acts as a mild regulariser.

WHY 3 CONV LAYERS?
Layer 1 detects edges/colour transitions; Layer 2 assembles block boundaries; Layer 3 captures full cluster patterns.

WHY DROPOUT(0.4)?
Applied to dense head only — prevents co-adaptation of 64 learned features without disrupting convolutional feature extraction.

INPUT — 30-DAY ROLLING CORRELATION HEATMAPS

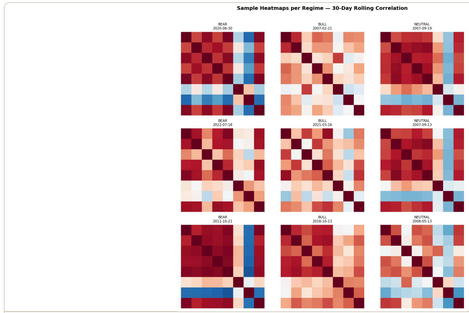

Bear: uniform red equity block + blue TLT/GLD columns — crisis correlation spike. Bull: muted, mixed structure. Neutral: intermediate. Fixed vmin=−1, vmax=+1 encodes absolute magnitude as pixel intensity, enabling direct regime-to-regime comparison.

CLASS IMBALANCE — PROBLEM → FIX → RESIDUAL EFFECT

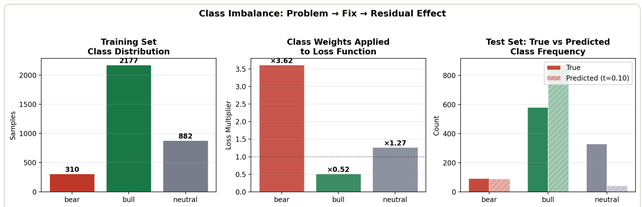

Training set is 65% bull. Class weights (bear ×3.62, neutral ×1.27, bull ×0.52) scale the cross-entropy loss so rare-class errors penalise the network more. Right panel: despite the fix, the model still over-predicts bull on the test set, reflecting the bear class's genuinely ambiguous visual signature outside crisis periods.

THRESHOLD ANALYSIS — PER-CLASS F1 VS BEAR THRESHOLD

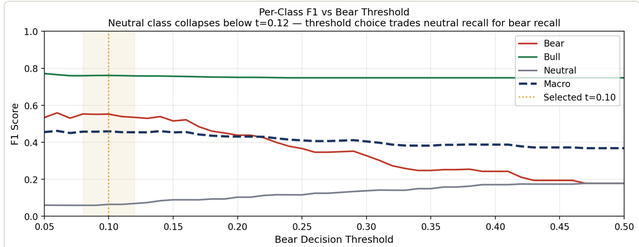

Bear F1 peaks near t=0.05–0.12 then declines steeply. Neutral collapses below t=0.12 — the bear threshold directly controls how much of the ambiguous neutral space is absorbed into bear predictions. t=0.10 is chosen as the operating point maximising bear recall while keeping macro-F1 above the argmax baseline.

REGIME TIMELINE — TEST PERIOD 2021-2024

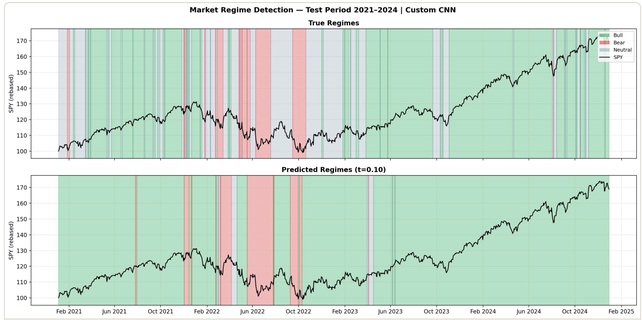

Top panel: ground-truth regimes shaded over rebased SPY. Bottom: model predictions at t=0.10. The 2022 bear cluster (Jun–Oct, −25% drawdown) is captured correctly. Most 2021 bear days are detected. The neutral class is near-absent in predictions — the low threshold converts borderline neutral days into bear calls.

TRAINING CURVES — CUSTOM CNN VS MOBILENETV2

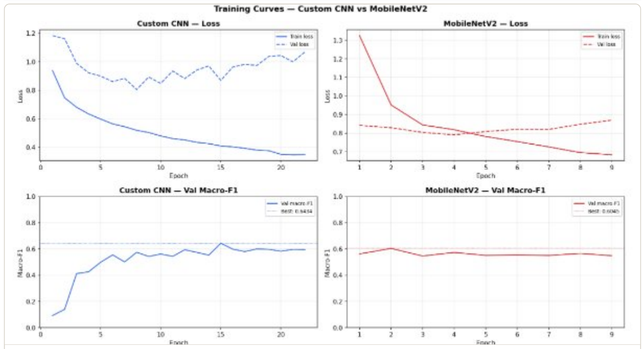

Custom CNN: loss converges steadily; val macro-F1 peaks at 0.643 (epoch 15), early-stopped at epoch 22. MobileNetV2 (frozen ImageNet backbone): val macro-F1 plateaus at 0.605 by epoch 2 and stagnates — photographic low-level features do not transfer to numerical colour-encoded correlation matrices.

MODEL VS BASELINE COMPARISON

ABLATION — WHAT EACH CHOICE CONTRIBUTES

Domain CNN learns correlation block structure — ImageNet features miss this entirely (+14 F1 vs MobileNetV2)

Threshold t=0.10 recovers 51 pts of bear recall vs argmax without any retraining (+5 macro-F1 pts)

Class weights force bear learning; without them bear recall ≈ 0% (model predicts bull for everything)

Fixed colour scale (vmin=−1, vmax=+1) is essential — adaptive scaling would erase the regime signal

CONFUSION MATRIX (T = 0.10, N = 1,004 TEST DAYS)

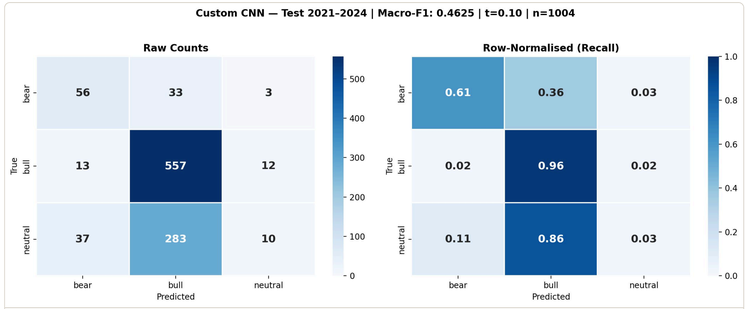

Left: raw counts. Right: row-normalised recall. Bear recall 61% (56/92), Bull 96%, Neutral 3%. The model overwhelmingly defaults uncertain inputs to bull — neutral has no sharp visual signature that separates it from low-volatility bull periods.

CLASSIFICATION REPORT — TEST SET

Class	Precision	Recall	F1	Support
Bear	0.528	0.609	0.566	92
Bull	0.638	0.957	0.766	582
Neutral	0.400	0.030	0.056	330
Macro avg	**0.522**	**0.532**	**0.463**	**1,004**

HEAD-TO-HEAD

Model	Val Macro-F1	Test Macro-F1
Custom CNN (argmax)	0.643	0.414
Custom CNN t=0.10 ★	**—**	**0.463**
MobileNetV2 (frozen)	0.605	~0.32

102k-param domain-specific CNN outperforms 2.3M-param ImageNet backbone by +14 macro-F1 pts. Smaller, faster, and purpose-built.

KEY FINDING

Bear recall: 9.8% → 61% via threshold calibration at t=0.10 — no retraining required. The network had learned a genuine bear signal; it was simply miscalibrated by the class imbalance. Lowering the decision boundary surfaces it. This is the central result: **architecture was right, calibration was wrong.**

GRAD-CAM — MODEL ATTENTION MAPS

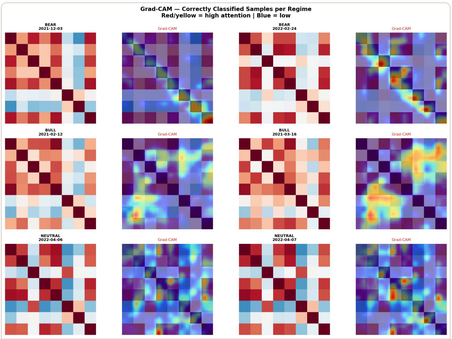

*Two correctly-classified samples per regime. **Bear:** high attention on off-diagonal cells where TLT/GLD decouple — confirms the model learned genuine crisis structure. **Bull:** diffuse, low-magnitude. **Neutral:** scattered — no dominant correlation feature to attend to.*

PRECISION-RECALL CURVES — MODEL LIFT OVER RANDOM BASELINE

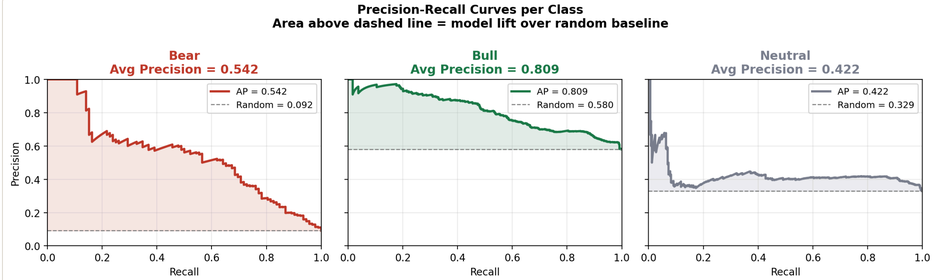

*Shaded area = lift above random (dashed). **Bear AP = 0.542** vs random = 0.092 → **5.9× lift**. Bull AP = 0.809. Neutral AP = 0.422 barely clears random (0.329) — confirming weak separability in correlation space for the neutral class.*

KEY FINDINGS

9.8%→61%
Bear recall via threshold calibration at t=0.10 — no retraining required

5.9× lift
Bear AP over **random baseline** — genuine learned signal confirmed

+14 F1
Custom CNN over **MobileNetV2** — domain architecture beats ImageNet priors

3% recall
Neutral class — no sharp visual signature; AP only 0.09 above random

WHY CNN FOR THIS TASK?

Correlation matrices are **structured 2D inputs** — equity assets cluster top-left, macro instruments (TLT, GLD) separate bottom-right. A CNN's **inductive bias** (local receptive fields, translation equivariance) directly matches this: a 3×3 kernel detects a local correlation transition; stacked layers compose these into block-level regime patterns.

Alternatives (MLP, tree methods) require **flattening** the matrix, destroying the 2D spatial structure that carries the signal. Grad-CAM confirms the network exploits spatial structure — attending to specific off-diagonal regions, *not* the uninformative diagonal (which is always +1).

LIMITATIONS & FUTURE WORK

Neutral class: residual labelling gives no clean visual identity. A richer labelling scheme (late-bull / early-bear sub-regimes) would likely improve recall significantly.

Calibration gap: model trained on COVID-style crises (val 2019–20); 2022 slow structural bear required threshold recalibration. A multi-period validation set would yield better-calibrated probabilities.

Future directions: eigenvalue decomposition as side-channel · temporal smoothing prior · fine-tuning MobileNetV2 on domain data · regime-conditioned trading strategy overlay.

REPRODUCIBILITY

`RUN_FULL = False` — load saved artifacts, end-to-end in ~30 sec on Kaggle T4

`RUN_FULL = True` — download prices, generate ~4,900 PNGs, train both models, ~25 min T4

Attach dataset: `market-regime-cnn-artifacts` · *Companion:* `market_regime_cnn.ipynb`